NAME

Encode - character encodings

SYNOPSIS

```
use Encode;
```

Table of Contents

Encode consists of a collection of modules whose details are too big to fit in one document. This POD itself explains the top-level APIs and general topics at a glance. For other topics and more details, see the PODs below:

```
Name             Description
-----------------------------------------------------
Encode::Alias        Alias definitions to encodings
Encode::Encoding     Encode Implementation Base Class
Encode::Supported    List of Supported Encodings
Encode::CN           Simplified Chinese Encodings
Encode::JP           Japanese Encodings
Encode::KR           Korean Encodings
Encode::TW           Traditional Chinese Encodings
-----------------------------------------------------
```

DESCRIPTION

The `Encode` module provides the interfaces between Perl's strings and the rest of the system. Perl strings are sequences of **characters**.

The repertoire of characters that Perl can represent is at least that defined by the Unicode Consortium. On most platforms the ordinal values of the characters (as returned by `ord(ch)`) is the "Unicode codepoint" for the character (the exceptions are those platforms where the legacy encoding is some variant of EBCDIC rather than a super-set of ASCII - see *perlebcdic*).

Traditionally, computer data has been moved around in 8-bit chunks often called "bytes". These chunks are also known as "octets" in networking standards. Perl is widely used to manipulate data of many types - not only strings of characters representing human or computer languages but also "binary" data being the machine's representation of numbers, pixels in an image - or just about anything.

When Perl is processing "binary data", the programmer wants Perl to process "sequences of bytes". This is not a problem for Perl - as a byte has 256 possible values, it easily fits in Perl's much larger "logical character".

TERMINOLOGY

- *character*: a character in the range 0..(2**32-1) (or more). (What Perl's strings are made of.)

- *byte*: a character in the range 0..255 (A special case of a Perl character.)

- *octet*: 8 bits of data, with ordinal values 0..255 (Term for bytes passed to or from a non-Perl context, e.g. a disk file.)

PERL ENCODING API

$octets = encode(ENCODING, $string [, CHECK])

Encodes a string from Perl's internal form into *ENCODING* and returns a sequence of octets. ENCODING can be either a canonical name or an alias. For encoding names and aliases, see *Defining Aliases*. For CHECK, see *Handling Malformed Data*.

For example, to convert a string from Perl's internal format to iso-8859-1 (also known as Latin1),

```
$octets = encode("iso-8859-1", $string);
```

CAVEAT: When you run `$octets = encode("utf8", $string)`, then $octets **may not be equal to** $string. Though they both contain the same data, the UTF8 flag for $octets is **always** off. When you encode anything, UTF8 flag of the result is always off, even when it contains completely valid utf8 string. See *The UTF8 flag* below.

If the $string is `undef` then `undef` is returned.

$string = decode(ENCODING, $octets [, CHECK])

Decodes a sequence of octets assumed to be in *ENCODING* into Perl's internal form and returns the resulting string. As in encode(), ENCODING can be either a canonical name or an alias. For encoding names and aliases, see *Defining Aliases*. For CHECK, see *Handling Malformed Data*.

For example, to convert ISO-8859-1 data to a string in Perl's internal format:

```
$string = decode("iso-8859-1", $octets);
```

CAVEAT: When you run `$string = decode("utf8", $octets)`, then $string **may not be equal to** $octets. Though they both contain the same data, the UTF8 flag for $string is on unless $octets entirely consists of ASCII data (or EBCDIC on EBCDIC machines). See *The UTF8 flag* below.

If the $string is `undef` then `undef` is returned.

[$obj =] find_encoding(ENCODING)

Returns the *encoding object* corresponding to ENCODING. Returns undef if no matching ENCODING is find.

This object is what actually does the actual (en|de)coding.

```
$utf8 = decode($name, $bytes);
```

is in fact

```
$utf8 = do{
  $obj = find_encoding($name);
  croak qq(encoding "$name" not found) unless ref $obj;
  $obj->decode($bytes)
};
```

with more error checking.

Therefore you can save time by reusing this object as follows;

```
my $enc = find_encoding("iso-8859-1");
while(<>){
    my $utf8 = $enc->decode($_);
    # and do someting with $utf8;
}
```

Besides `->decode` and `->encode`, other methods are available as well. For instance, `-> name` returns the canonical name of the encoding object.

```
find_encoding("latin1")->name; # iso-8859-1
```

See *Encode::Encoding* for details.

[$length =] from_to($octets, FROM_ENC, TO_ENC [, CHECK])

Converts **in-place** data between two encodings. The data in $octets must be encoded as octets and not as characters in Perl's internal format. For example, to convert ISO-8859-1 data to Microsoft's CP1250 encoding:

```
from_to($octets, "iso-8859-1", "cp1250");
```

and to convert it back:

```
from_to($octets, "cp1250", "iso-8859-1");
```

Note that because the conversion happens in place, the data to be converted cannot be a string constant; it must be a scalar variable.

from_to() returns the length of the converted string in octets on success, *undef* on error.

CAVEAT: The following operations look the same but are not quite so;

```
from_to($data, "iso-8859-1", "utf8"); #1
$data = decode("iso-8859-1", $data);  #2
```

Both #1 and #2 make $data consist of a completely valid UTF-8 string but only #2 turns UTF8 flag on. #1 is equivalent to

```
$data = encode("utf8", decode("iso-8859-1", $data));
```

See *The UTF8 flag* below.

Also note that

```
from_to($octets, $from, $to, $check);
```

is equivalent to

```
$octets = encode($to, decode($from, $octets), $check);
```

Yes, it does not respect the $check during decoding. It is deliberately done that way. If you need minute control, `decode` then `encode` as follows;

```
$octets = encode($to, decode($from, $octets, $check_from), $check_to);
```

$octets = encode_utf8($string);

Equivalent to `$octets = encode("utf8", $string);` The characters that comprise $string are encoded in Perl's internal format and the result is returned as a sequence of octets. All possible characters have a UTF-8 representation so this function cannot fail.

$string = decode_utf8($octets [, CHECK]);

equivalent to `$string = decode("utf8", $octets [, CHECK])`. The sequence of octets represented by $octets is decoded from UTF-8 into a sequence of logical characters. Not all sequences of octets form valid UTF-8 encodings, so it is possible for this call to fail. For CHECK, see *Handling Malformed Data*.

Listing available encodings

```
use Encode;
@list = Encode->encodings();
```

Returns a list of the canonical names of the available encodings that are loaded. To get a list of all available encodings including the ones that are not loaded yet, say

```
@all_encodings = Encode->encodings(":all");
```

Or you can give the name of a specific module.

```
@with_jp = Encode->encodings("Encode::JP");
```

When "::" is not in the name, "Encode::" is assumed.

```
@ebcdic = Encode->encodings("EBCDIC");
```

To find out in detail which encodings are supported by this package, see *Encode::Supported*.

Defining Aliases

To add a new alias to a given encoding, use:

```
use Encode;
use Encode::Alias;
define_alias(newName => ENCODING);
```

After that, newName can be used as an alias for ENCODING. ENCODING may be either the name of an encoding or an *encoding object*

But before you do so, make sure the alias is nonexistent with `resolve_alias()`, which returns the canonical name thereof. i.e.

```
Encode::resolve_alias("latin1") eq "iso-8859-1" # true
Encode::resolve_alias("iso-8859-12")   # false; nonexistent
Encode::resolve_alias($name) eq $name  # true if $name is canonical
```

resolve_alias() does not need `use Encode::Alias`; it can be exported via `use Encode qw(resolve_alias)`.

See *Encode::Alias* for details.

Finding IANA Character Set Registry names

The canonical name of a given encoding does not necessarily agree with IANA IANA Character Set Registry, commonly seen as `Content-Type: text/plain; charset=`*whatever*. For most cases canonical names work but sometimes it does not (notably 'utf-8-strict').

Therefore as of Encode version 2.21, a new method `mime_name()` is added.

```
use Encode;
my $enc = find_encoding('UTF-8');
warn $enc->name;      # utf-8-strict
warn $enc->mime_name; # UTF-8
```

See also: *Encode::Encoding*

Encoding via PerlIO

If your perl supports *PerlIO* (which is the default), you can use a PerlIO layer to decode and encode directly via a filehandle. The following two examples are totally identical in their functionality.

```
# via PerlIO
open my $in,  "<:encoding(shiftjis)", $infile  or die;
open my $out, ">:encoding(euc-jp)",   $outfile or die;
while(<$in>){ print $out $_; }

# via from_to
open my $in,  "<", $infile  or die;
open my $out, ">", $outfile or die;
while(<$in>){
  from_to($_, "shiftjis", "euc-jp", 1);
  print $out $_;
}
```

Unfortunately, it may be that encodings are PerlIO-savvy. You can check if your encoding is supported by PerlIO by calling the `perlio_ok` method.

```
Encode::perlio_ok("hz");              # False
find_encoding("euc-cn")->perlio_ok;  # True where PerlIO is available


use Encode qw(perlio_ok);            # exported upon request
perlio_ok("euc-jp")
```

Fortunately, all encodings that come with Encode core are PerlIO-savvy except for hz and ISO-2022-kr. For gory details, see *Encode::Encoding* and *Encode::PerlIO*.

Handling Malformed Data

The optional *CHECK* argument tells Encode what to do when it encounters malformed data. Without CHECK, Encode::FB_DEFAULT (== 0) is assumed.

As of version 2.12 Encode supports coderef values for CHECK. See below.

NOTE: Not all encoding support this feature

Some encodings ignore *CHECK* argument. For example, *Encode::Unicode* ignores *CHECK* and it always croaks on error.

Now here is the list of *CHECK* values available

CHECK = Encode::FB_DEFAULT (== 0)

If *CHECK* is 0, (en|de)code will put a *substitution character* in place of a malformed character. When you encode, <subchar> will be used. When you decode the code point $0xFFFD$ is used. If the data is supposed to be UTF-8, an optional lexical warning (category utf8) is given.

CHECK = Encode::FB_CROAK (== 1)

If *CHECK* is 1, methods will die on error immediately with an error message. Therefore, when *CHECK* is set to 1, you should trap the error with eval{} unless you really want to let it die.

CHECK = Encode::FB_QUIET

If *CHECK* is set to Encode::FB_QUIET, (en|de)code will immediately return the portion of the data that has been processed so far when an error occurs. The data argument will be overwritten with everything after that point (that is, the unprocessed part of data). This is handy when you have to call decode repeatedly in the case where your source data may contain partial multi-byte character sequences, (i.e. you are reading with a fixed-width buffer). Here is a sample code that does exactly this:

```
my $buffer = ''; my $string = '';
while(read $fh, $buffer, 256, length($buffer)){
  $string .= decode($encoding, $buffer, Encode::FB_QUIET);
  # $buffer now contains the unprocessed partial character
}
```

CHECK = Encode::FB_WARN

This is the same as above, except that it warns on error. Handy when you are debugging the mode above.

perlqq mode (*CHECK* = Encode::FB_PERLQQ)

HTML charref mode (*CHECK* = Encode::FB_HTMLCREF)

XML charref mode (*CHECK* = Encode::FB_XMLCREF)

For encodings that are implemented by Encode::XS, CHECK == Encode::FB_PERLQQ turns (en|de)code into `perlqq` fallback mode.

When you decode, $\x{HH}$ will be inserted for a malformed character, where *HH* is the hex representation of the octet that could not be decoded to utf8. And when you encode, $\x\{HHHH\}$ will be inserted, where *HHHH* is the Unicode ID of the character that cannot be found in the character repertoire of the encoding.

HTML/XML character reference modes are about the same, in place of $\x\{HHHH\}$, HTML uses `&#NNN;` where *NNN* is a decimal number and XML uses `&#xHHHH;` where *HHHH* is the hexadecimal number.

In Encode 2.10 or later, `LEAVE_SRC` is also implied.

The bitmask

These modes are actually set via a bitmask. Here is how the FB_XX constants are laid out. You can import the FB_XX constants via `use Encode qw(:fallbacks);` you can import the generic bitmask constants via `use Encode qw(:fallback_all)`.

```
                      FB_DEFAULT FB_CROAK FB_QUIET FB_WARN  FB_PERLQQ
DIE_ON_ERR    0x0001              X
WARN_ON_ERR   0x0002                                 X
RETURN_ON_ERR 0x0004                       X         X
LEAVE_SRC     0x0008                                          X
PERLQQ        0x0100                                          X
HTMLCREF      0x0200
XMLCREF       0x0400
```

Encode::LEAVE_SRC

If the `Encode::LEAVE_SRC` bit is not set, but *CHECK* is, then the second argument to `encode()` or `decode()` may be assigned to by the functions. If you're not interested in this, then bitwise-or the bitmask with it.

coderef for CHECK

As of Encode 2.12 CHECK can also be a code reference which takes the ord value of unmapped caharacter as an argument and returns a string that represents the fallback character. For instance,

```
$ascii = encode("ascii", $utf8, sub{ sprintf "<U+%04X>", shift });
```

Acts like FB_PERLQQ but <U+*XXXX*> is used instead of \x{*XXXX*}.

Defining Encodings

To define a new encoding, use:

```
use Encode qw(define_encoding);
define_encoding($object, 'canonicalName' [, alias...]);
```

canonicalName will be associated with *$object*. The object should provide the interface described in *Encode::Encoding*. If more than two arguments are provided then additional arguments are taken as aliases for *$object*.

See *Encode::Encoding* for more details.

The UTF8 flag

Before the introduction of Unicode support in perl, The `eq` operator just compared the strings represented by two scalars. Beginning with perl 5.8, `eq` compares two strings with simultaneous consideration of *the UTF8 flag*. To explain why we made it so, I will quote page 402 of `Programming Perl, 3rd ed.`

Goal #1:

Old byte-oriented programs should not spontaneously break on the old byte-oriented data they

used to work on.

Goal #2:

Old byte-oriented programs should magically start working on the new character-oriented data when appropriate.

Goal #3:

Programs should run just as fast in the new character-oriented mode as in the old byte-oriented mode.

Goal #4:

Perl should remain one language, rather than forking into a byte-oriented Perl and a character-oriented Perl.

Back when `Programming Perl, 3rd ed.` was written, not even Perl 5.6.0 was born and many features documented in the book remained unimplemented for a long time. Perl 5.8 corrected this and the introduction of the UTF8 flag is one of them. You can think of this perl notion as of a byte-oriented mode (UTF8 flag off) and a character-oriented mode (UTF8 flag on).

Here is how Encode takes care of the UTF8 flag.

- When you encode, the resulting UTF8 flag is always off.

- When you decode, the resulting UTF8 flag is on unless you can unambiguously represent data. Here is the definition of dis-ambiguity.

  After `$utf8 = decode('foo', $octet);`,

  ```
  When $octet is...   The UTF8 flag in $utf8 is
  ----------------------------------------------
  In ASCII only (or EBCDIC only)          OFF
  In ISO-8859-1                            ON
  In any other Encoding                    ON
  ----------------------------------------------
  ```

  As you see, there is one exception, In ASCII. That way you can assume Goal #1. And with Encode Goal #2 is assumed but you still have to be careful in such cases mentioned in **CAVEAT** paragraphs.

  This UTF8 flag is not visible in perl scripts, exactly for the same reason you cannot (or you *don't have to*) see if a scalar contains a string, integer, or floating point number. But you can still peek and poke these if you will. See the section below.

Messing with Perl's Internals

The following API uses parts of Perl's internals in the current implementation. As such, they are efficient but may change.

is_utf8(STRING [, CHECK])

[INTERNAL] Tests whether the UTF8 flag is turned on in the STRING. If CHECK is true, also checks the data in STRING for being well-formed UTF-8. Returns true if successful, false otherwise.

As of perl 5.8.1, *utf8* also has utf8::is_utf8().

_utf8_on(STRING)

[INTERNAL] Turns on the UTF8 flag in STRING. The data in STRING is **not** checked for being well-formed UTF-8. Do not use unless you **know** that the STRING is well-formed UTF-8. Returns the previous state of the UTF8 flag (so please don't treat the return value as indicating success or failure), or `undef` if STRING is not a string.

This function does not work on tainted values.

_utf8_off(STRING)

> [INTERNAL] Turns off the UTF8 flag in STRING. Do not use frivolously. Returns the previous state of the UTF8 flag (so please don't treat the return value as indicating success or failure), or `undef` if STRING is not a string.

> This function does not work on tainted values.

UTF-8 vs. utf8 vs. UTF8

```
....We now view strings not as sequences of bytes, but as sequences
of numbers in the range 0 .. 2**32-1 (or in the case of 64-bit
computers, 0 .. 2**64-1) -- Programming Perl, 3rd ed.
```

That has been the perl's notion of UTF-8 but official UTF-8 is more strict; Its ranges is much narrower (0 .. 10FFFF), some sequences are not allowed (i.e. Those used in the surrogate pair, 0xFFFE, et al).

Now that is overruled by Larry Wall himself.

```
From: Larry Wall <larry@wall.org>
Date: December 04, 2004 11:51:58 JST
To: perl-unicode@perl.org
Subject: Re: Make Encode.pm support the real UTF-8
Message-Id: <20041204025158.GA28754@wall.org>


On Fri, Dec 03, 2004 at 10:12:12PM +0000, Tim Bunce wrote:
: I've no problem with 'utf8' being perl's unrestricted uft8 encoding,
: but "UTF-8" is the name of the standard and should give the
: corresponding behaviour.

For what it's worth, that's how I've always kept them straight in my
head.

Also for what it's worth, Perl 6 will mostly default to strict but
make it easy to switch back to lax.

Larry
```

Do you copy? As of Perl 5.8.7, **UTF-8** means strict, official UTF-8 while **utf8** means liberal, lax, version thereof. And Encode version 2.10 or later thus groks the difference between `UTF-8` and C"utf8".

```
encode("utf8",  "\x{FFFF_FFFF}", 1); # okay
encode("UTF-8", "\x{FFFF_FFFF}", 1); # croaks
```

`UTF-8` in Encode is actually a canonical name for `utf-8-strict`. Yes, the hyphen between "UTF" and "8" is important. Without it Encode goes "liberal"

```
find_encoding("UTF-8")->name # is 'utf-8-strict'
find_encoding("utf-8")->name # ditto. names are case insensitive
find_encoding("utf_8")->name  # ditto. "_" are treated as "-"
find_encoding("UTF8")->name  # is 'utf8'.
```

The UTF8 flag is internally called UTF8, without a hyphen. It indicates whether a string is internally encoded as utf8, also without a hypen.

SEE ALSO

Encode::Encoding, Encode::Supported, Encode::PerlIO, encoding, perlebcdic, "open" in perlfunc, perlunicode, perluniintro, perlunifaq, perlunitut utf8, the Perl Unicode Mailing List < perl-unicode@perl.org>

MAINTAINER

This project was originated by Nick Ing-Simmons and later maintained by Dan Kogai < dankogai@dan.co.jp>. See AUTHORS for a full list of people involved. For any questions, use < perl-unicode@perl.org> so we can all share.

While Dan Kogai retains the copyright as a maintainer, the credit should go to all those involoved. See AUTHORS for those submitted codes.

COPYRIGHT

Copyright 2002-2006 Dan Kogai <dankogai@dan.co.jp>

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself.